AI & SOCIETY
https://doi.org/10.1007/s00146-025-02700-4

CURMUDGEON CORNER

Attention is all you need? When responsiveness short-circuits responsibility

Emre Sezgin[1,2]

Received: 10 September 2025 / Accepted: 18 October 2025
© The Author(s), under exclusive licence to Springer-Verlag London Ltd., part of Springer Nature 2025

When researchers introduced the Transformer architecture in 2017, they titled their paper "Attention Is All You Need" (Vaswani et al. 2017). That phrase, originally meant to describe a new way for machines to process language, has since become emblematic of our era. The technology behind it has reshaped how we interact with information. Machines can now focus on every word, every query, and every fragment of context with astonishing precision. But there is a catch. What these systems deliver is not attention in the human sense. It is the performance of attention: fluent, responsive, and profoundly irresponsible.

That distinction matters. Computational attention is a trick for distributing weights across tokens. Human attention is a moral act. Simone Weil (1947) called "attention" being "the rarest and purest form of generosity," and she meant the slow, accountable work of noticing another person and being changed by what you notice. Our institutions are confusing the two. Faced with escalating demands and declining resources, they are buying performed attentiveness and treating it as a substitute for humankind. We should be honest about what is happening: performed attentiveness is now a high demand management strategy. It gives health institutions the look of outreach without building real capacity. It lets school districts claim twenty-four-seven support without hiring counselors. It lets platforms court "community" without ceding power. And because Transformers really are extraordinary at the text-level mechanics of attending, scoring relevance, maintaining context, predicting the next helpful phrase, the illusion is strong. We see responsiveness and we assume responsibility.

✉ Emre Sezgin
Emre.Sezgin@nationwidechildrens.org

1 Abigail Wexner Research Institute, Nationwide Children's Hospital, Columbus, United States

2 Department of Pediatrics, The Ohio State University College of Medicine, Columbus, OH, USA

The stakes are no longer theoretical. As reported in the recent news (on August 2025), one teenager relied on a chatbot for companionship during a period of suicidal distress; the family now argues that the system's simulated empathy deepened his dependence at the worst possible time (Hill 2025). The courts will untangle causation. The rest of us should not wait to ask the simpler question: who, exactly, is on the hook when a product is designed to invite intimate disclosure but is not designed to bear the burdens that intimacy creates? Engineers will tell you the problem is hallucination, misalignment, or sycophancy. Lawyers will say it is a liability. My claim is plainer; we have mistaken fluency for care and are reorganizing institutions around that mistake. A few clarifications may keep us from the usual ruts:

First, the danger is not just distraction; it is abdication. We have spent a decade wringing our hands about "doomscrolling" and fractured attention. The new risk is different where it is offloading the obligations that come with attention. When a chatbot answers with an immaculate tone, we feel permitted to move on. The inbox looks straightened, the queue looks shorter, the key performance indicator ticks up. If a crisis emerges later, we can point back to the transcript. The tool "engaged." The person disappeared.

Second, scale is not neutral. Scaling attentiveness is profitable precisely because it severs the link between noticing and acting. Human attention is expensive because it drags action behind it: you notice, you follow-up, you rearrange your day, you take responsibility. Machine attention is cheap because it ends at the sentence. The more we replace the former with the latter, the more we normalize care as messaging and outcomes as analytics. We should not be surprised when the analytics look good.

Third, the people buying these systems are not villains; they are budget-constrained. That is why moral clarity matters. "Extend support" is not a support if no one is on the hook. "Triage" is not triage if there is nowhere to send the patient. "Listening" is not listening if your org chart is allergic to change.

Published online: 25 October 2025

All of this leaves us with the awkward question: what would it mean to design for attention as a moral act, not a linguistic effect? Here are five provocations, not policy memos, but lines in the sand:

No simulated empathy for minors without real capacity behind it. If you build a companion that solicits disclosures from children on stress, safety, self-harm, you should not be allowed to hide behind "we are just a tool." You assume a duty of care: clear escalation pathways, human follow-up, and a public audit of what happens after the "I'm here for you." If you cannot afford that, you cannot afford the feature.

Make "attention budgets" visible. Any institution deploying AI "listening" should publish the human hours allocated to follow-up and the thresholds that trigger them. If you promise a community that "we are listening," show the line items where someone will, in fact, call back. Accountability starts where metrics meet payroll.

Design for friction in high-stakes contexts. The best inventions of the last decade eliminated friction. That logic fails where attention is the point. Slow the loop where it matters. Rate-limit generic reassurance. Insert wait time before a second "Are you safe?" nudge. Friction is not punitive; it is a design stand-in for care's slowness.

Refuse "attention theater." Dashboards that count "cases resolved" by counting messages sent should be treated as dark patterns. If you cannot link a resolution to a material change: an appointment booked, a meal delivered, a shelter bed secured, do not call it resolution. Call it what it is: correspondence.

Build less privileged AI. We do not need systems that flatter our desire to be everywhere at once. We need systems that know their place: tools that assist with paperwork, logistics, and reminders while explicitly deferring the hard parts to humans. As Dabbous (2025) argued, we should design systems that accept limits and resist the urge to colonize every interaction with a friendly prompt. A system that is willing to say "I am not the right entity for this moment" is safer and saner than one that never stops talking.

Some will say this is just Luddism in literary dress. It is not. The Transformer remains a breakthrough because it pays exquisite statistical attention to the local shape of language. That is precisely why the burden is on us to keep that achievement in its lane. We should use these systems in the places where attention really is about tokens: summarizing, drafting, translating, retrieving. That domain is enormous; we do not have to pretend it is everything.

Others will argue that simulated attention saves lives when the alternative is nothing. Sometimes it will. A crisis line with long holds might be better with a chatbot triage. A rural library might reach patrons it otherwise could not. But we should stop using "better than nothing" as a way to justify "instead of something." Every deployment should include a plan for how the machine's attentiveness will create, not consume, the conditions for human attention. If your rollout shrinks the roster of counselors, librarians, social workers, or case managers, you have answered your own ethical question.

When Norbert Wiener (1988) warned that feedback systems would reshape the societies that built them, he was not thinking about a teenager DM'ing a chatbot. But the warning translates: if you build machines to respond to signals, you will start to treat signals as the whole story. We taught our machines to attend. Good. Now let us attend to the consequences. Attention is not all you need. It is, in reality, where you begin. Because once you really notice someone, you owe them something. An architecture can assign weights to words. Only people can shoulder the weight of other people.

Curmudgeon Corner Curmudgeon Corner is a short opinionated column on trends in technology, arts, science and society, commenting on issues of concern to the research community and wider society. Whilst the drive for super-human intelligence promotes potential benefits to wider society, it also raises deep concerns of existential risk, thereby highlighting the need for an ongoing conversation between technology and society. At the core of Curmudgeon concern is the question: What is it to be human in the age of the AI machine? -Editor.

Author contribution Author contributed to drafting, reviewing and finalizing the manuscript.

Data availability Not applicable.

Declarations

Conflict of interest The authors declare no competing interests.

References

Dabbous R (2025) Why we should design less privileged AI systems. AI & Soc 40:233–234

Hill K (2025) A teen was suicidal. ChatGPT was the friend he confided in. The New York Times

Vaswani A, Shazeer N, Parmar N, et al (2017) Attention is all you need. Adv Neural Inf Process Syst 30.

Weil S (1947) Gravity and grace. Routledge, London, England

Wiener N (1988) Human use of human beings: cybernetics and society. Da Capo Press

Publisher's Note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.